

Our range of enforcement options

When we take enforcement actions, we may do so either on a specific piece of content (e.g., an individual Tweet or Direct Message) or on an ~~account~~. We may employ a combination of these options. In some instances, this is because the behavior violates the [Twitter Rules](<http://web.archive.org/en/rules-and-policies/twitter-rules.html>). Other times, it may be in response to a valid ~~and properly-sourced~~ request from an authorized entity in a given country. ~~Below are some of~~ the enforcement actions that we may take.

Tweet-level enforcement

We take action at the Tweet level to ensure that we are not being overly harsh with an otherwise healthy account that made a mistake and violated our rules.

A few of the ways in which we might take action at the Tweet level include:

~~Labeling a Tweet that may contain disputed or misleading information. Our current misleading information policies cover synthetic and manipulated media~~(<http://web.archive.org/en/rules-and-policies/manipulated-media.html>) ([CIVIC ID](https://blog.twitter.com/en-us/topics/company/2022/civic-integrity-notices))(<https://blog.twitter.com/en-us/topics/company/2022/civic-integrity-notices>) and ~~(civic integrity)~~(<http://web.archive.org/en/rules-and-policies/election-integrity-policy.html>). If we determine a Tweet contains misleading or disputed information per our policies that could lead to harm, we may ~~add a label~~(<http://web.archive.org/en/rules-and-policies/notices-on-twitter.html>) to the content to provide context and additional information.

~~Limiting Tweet visibility. This makes content less visible on Twitter, either by making Tweets ineligible for amplification in Top search results and on timelines for users who don't follow the Tweet author, by downranking Tweets in replies (except when the user follows the Tweet author), and/or excluding Tweets and/or accounts in email or in-product recommendations. Limiting Tweet visibility depends on a number of signals about the nature of the interaction and the type of the content.~~

Requiring Tweet removal: When we determine that a Tweet violated the Twitter Rules, we ~~require the violator to remove it before they can Tweet again. We send an email notification to the violator identifying the Tweet(s) in violation and which policies have been violated. They will then need to go through the process of removing the violating Tweet or (appealing)~~(<https://help.twitter.com/forms/general?subtopic=suspended>) our review if they believe we made an error.

~~Hiding a violating Tweet while awaiting its removal. In the interim period between when Twitter takes enforcement action and the person removes the Tweet, we hide that Tweet from public view and will replace the original content with a notice stating that the Tweet is no longer available because it violated our Rules. Additionally, this notice will be available for 14 days after the Tweet is removed.~~

Where this notice is available

- ~~The details page of the Tweet in violation on Twitter for iOS and Android, and [Twitter.com](#).~~
- ~~The profile page of the account in violation on Twitter for iOS and Android, and [Twitter.com](#).~~

Where this notice is not available

- ~~Areas outside of profiles and Tweet details pages, such as Home timeline, notifications, and search of any Twitter client.~~

Notice of public interest exception: ~~In rare cases, we may determine that it is in the public interest for a Tweet that would otherwise be in violation of our rules to remain accessible on our service. [Learn more]~~(<http://web.archive.org/en/rules-and-policies/public-interest.html>) about the public interest exception.

When we make such an exception, we'll place the Tweet behind a notice explaining the exception and giving you the option to view the Tweet if you wish.

When applying the notice, we'll also take the following actions to reduce the Tweet's visibility on our service:

- ~~We'll turn off engagements like replies, Retweets, and likes, if you want to talk about it, you can still Quote Tweet.~~
- ~~We won't show any engagement counts on the Tweet (e.g. number of likes, replies), but any prior engagements can be undone (e.g. unliked, Retweets undone).~~
- ~~Any previous replies will not be viewable within the Tweet details.~~

Additionally, in order to minimize potential harm, Tweets that are placed behind the notice will not be available in the following areas of Twitter:

- ~~Tweets in the "Top Tweets" Home timeline~~
- ~~Safe search~~
- ~~Recommendations via push and Notifications tab~~
- ~~Email and text recommendations~~
- ~~Live event timeline~~
- ~~Explore tab~~

Direct Message-level enforcement

Stopping conversations between a reported violator and the reporter's

Our range of enforcement options

When we take enforcement actions, we may do so either on a specific piece of content (e.g., an individual Tweet or Direct Message), on an ~~account~~, or employ a combination of these options. In some instances, this is because the behavior violates the [Twitter Rules](<http://web.archive.org/en/rules-and-policies/twitter-rules.html>). Other times, it may be in response to a valid ~~legal~~ request from an authorized entity in a given country. ~~Below are some of~~ the enforcement actions that we may take.

Tweet-level enforcement

We take action at the Tweet level when a specific Tweet violates the Twitter Rules, including Tweets that share or reproduce other Tweets by posting screenshots, quote-Tweeting, or sharing Tweet URLs that violate our Rules.

A few of the ways in which we might take action at the Tweet level include:

Limiting Tweet visibility: Where appropriate, we will restrict the reach of Tweets that violate our policies and create a negative experience for other users by making the Tweet less discoverable on Twitter. This can include:

- Excluding the Tweet from search results, trends, and recommended notifications
- Removing the Tweet from the For you and Following timelines
- Restricting the Tweet's discoverability to the author's profile
- Downranking the Tweet in replies
- Restricting Likes, replies, Retweets, Quote Tweets, bookmarks, share, pin to profile, or Edit Tweet

Excluding the Tweet from having ads adjacent to it: Starting in April 2023, Tweets identified as violating our Rules will begin to receive labels informing both Tweet authors and viewers that we limited the Tweet's visibility.

Authors will be able to submit an appeal on the label if they think we incorrectly limited their Tweet's visibility.

Requiring Tweet removal: When we determine that a Tweet violated the Twitter Rules and the violation is severe enough to warrant Tweet removal, we will require the violator to remove it before they can Tweet again. They will need to go through the process of removing the violating Tweet or (appealing) (<http://web.archive.org/content/help-twitter/en/forms/account-access/appeals>) our removal request if they believe we made an error. The Tweet will be hidden from public view with a ~~(notice)~~(<http://web.archive.org/content/help-twitter/en/rules-and-policies/notices-on-twitter>) during this process.

Labeling a Tweet: If we determine a Tweet contains misleading or disputed information per our policies that could potentially lead to harm, we may ~~(add a label)~~(<http://web.archive.org/en/rules-and-policies/notices-on-twitter.html>) to the content to provide context and additional information to users. In these cases, ~~(Community Notes)~~(<http://web.archive.org/content/help-twitter/en/using-twitter/community-notes>) may also be visible on Tweets to provide additional context.

Notice of public interest exception: We may determine that it is in the public interest for a Tweet that would otherwise be in violation of our rules to remain accessible on our service. When this occurs, we will place the Tweet behind a notice and limit its visibility. ~~(Learn more)~~(<http://web.archive.org/en/rules-and-policies/public-interest.html>) about the public interest exception.

more)(<http://web.archive.org/en/rules-and-policies/public-interest.html>) about the public interest exception.

Direct Message-level enforcement

Stopping conversations between a reported violator and the reporter's account:

Our range of enforcement options

When we take enforcement actions, we may do so either on a specific piece of content (e.g., an individual ~~Tweet~~ or Direct Message), on an account, or ~~employ~~ a combination of these options. In some instances, this is because the behavior violates the ~~[Twitter Rules]~~(<http://web.archive.org/en/rules-and-policies/twitter-rules.html>). Other times, it may be in response to a valid ~~legal request from an authorized~~ entity in a given country. ~~Below are some of the enforcement actions that we~~ may take.

Tweet-level enforcement

We take action at the ~~Tweet~~ level when a specific ~~Tweet~~ violates the ~~Twitter Rules~~; including ~~Tweets~~ that share or reproduce other ~~Tweets~~ by posting ~~screenshots~~, ~~quote-Tweeting~~, or sharing ~~Tweet~~ URLs that violate our Rules.

A few of the ways in which we might take action at the ~~Tweet~~ level include:

Limiting ~~Tweet~~ visibility: Where appropriate, we will restrict the reach ~~of Tweets~~ that violate our policies and create a negative experience for other users by making the ~~Tweet~~ less discoverable on ~~Twitter~~. This can include:

- Excluding the ~~Tweet~~ from search results, trends, and recommended notifications
- Removing the ~~Tweet~~ from the For you and Following timelines
- Restricting the ~~Tweets~~ discoverability to the author's profile
- Downranking the ~~Tweet~~ in replies
- Restricting Likes, replies, ~~Retweets~~, Quote ~~Tweets~~, bookmarks, share, pin to profile, or Edit ~~Tweet~~

Excluding the ~~Tweet~~ from having ads adjacent to it: Starting in April ~~2023~~, ~~Tweets~~ identified as violating our Rules will begin to receive labels informing both ~~Tweet~~ authors and viewers that we limited the ~~Tweets~~ ~~visibility~~.

Authors will be able to submit an appeal on the label if they think we incorrectly limited their ~~Tweets~~ visibility.

Requiring ~~Tweet~~ removal: When we determine that a ~~Tweet~~ violated the ~~Twitter Rules~~ and the violation is severe enough to warrant ~~Tweet~~ removal, we ~~will require~~ the violator to remove it before they can ~~Tweet~~ again. They will ~~need to go~~ through the process of removing the violating ~~Tweet~~ or ~~(appealing)~~(<http://web.archive.org/content/help-twitter/en/forms/account-access/appeals>) our removal request if they believe we made an error. The ~~Tweet~~ will be hidden from public view with a ~~(notice)~~(<http://web.archive.org/content/help-twitter/en/rules-and-policies/notices-on-twitter>) during this process.

Labeling a ~~Tweet~~: If we determine a ~~Tweet~~ contains misleading or disputed information per our policies that could potentially lead to harm, we may ~~(add a label)~~(<http://web.archive.org/en/rules-and-policies/notices-on-twitter.html>) to the content to provide context and additional information to users. In these cases, ~~(Community Notes)~~(<http://web.archive.org/content/help-twitter/en/using-twitter/community-notes>) may also be visible on ~~Tweets~~ to provide additional context.

Notice of public interest exception: We may determine that it is in the public interest for a ~~Tweet~~ that would otherwise be in violation of our rules to remain accessible on our service. When this occurs, we will place the ~~Tweet~~ behind a notice and limit its visibility. ~~(Learn more)~~(<http://web.archive.org/en/rules-and-policies/public-interest.html>) about the public interest exception.

more)(<http://web.archive.org/en/rules-and-policies/public-interest.html>) about the public interest exception.

Direct Message-level enforcement

Stopping conversations between a reported violator and the reporter's account:

Our range of enforcement options

When we take enforcement actions, we may do so either on a specific piece of content (e.g., an individual ~~post~~ or Direct Message), on an account, or ~~employ~~ a combination of these options. In some instances, this is because the behavior violates the [X Rules]([/en/rules-and-policies/twitter-rules.html](https://en/rules-and-policies/twitter-rules.html)). Other times, it may be in response to a valid ~~legal request from an authorized~~ entity in a given country. ~~Below are some of the enforcement actions that we~~ may take.

Post-level enforcement

We take action at the ~~post~~ level when a specific ~~post~~ violates the X Rules, including posts that share or reproduce other posts by posting screenshots, quote-posting, or sharing post URLs that violate our Rules.

A few of the ways in which we might take action at the ~~post~~ level include:

Limiting ~~post~~ visibility: Where appropriate, we will restrict the reach ~~of posts~~ that violate our policies and create a negative experience for other users by making the ~~post~~ less discoverable on X. This can include:

- Excluding the ~~post~~ from search results, trends, and recommended notifications
- Removing the ~~post~~ from the For you and Following timelines
- Restricting the ~~post's~~ discoverability to the author's profile
- Downranking the ~~post~~ in replies
- Restricting Likes, replies, ~~Reposts~~, Quote ~~posts~~, bookmarks, share, pin to profile, or Edit ~~post~~

Excluding the ~~post~~ from having ads adjacent to it: Starting in April 2023, ~~posts~~ identified as violating our Rules will begin to receive labels informing both ~~post~~ authors and viewers that we limited the ~~post's~~ visibility.

Authors will be able to submit an appeal on the label if they think we incorrectly limited their ~~posts~~ visibility.

Requiring ~~post~~ removal: When we determine that a ~~post~~ violated the X Rules and the violation is severe enough to warrant ~~post~~ removal, we will ~~require~~ the violator to remove it before they can ~~post~~ again. They will ~~need to go~~ through the process of removing the violating ~~post~~ or (appealing) ([/en/forms/account-access/appeals.html](https://en/forms/account-access/appeals.html)) our removal request if they believe we made an error. The ~~post~~ will be hidden from public view with a ~~(notice)~~([/en/rules-and-policies/notices-on-x.html](https://en/rules-and-policies/notices-on-x.html)) during this process.

Labeling a ~~post~~: If we determine a ~~post~~ contains misleading or disputed information per our policies that could potentially lead to harm, we may ~~(add a label)~~([/en/rules-and-policies/notices-on-x.html](https://en/rules-and-policies/notices-on-x.html)) to the content to provide context and additional information to users. In these cases, ~~(Community Notes)~~([/en/using-twitter/community-notes.html](https://en/using-twitter/community-notes.html)) may also be visible on ~~posts~~ to provide additional context.

Notice of public interest exception: We may determine that it is in the public interest for a ~~post~~ that would otherwise be in violation of our rules to remain accessible on our service. When this occurs, we will place the ~~post~~ behind a notice and limit its visibility. ~~(Learn more)~~([/en/rules-and-policies/public-interest.html](https://en/rules-and-policies/public-interest.html)) about the public interest exception.

Direct Message-level enforcement

Stopping conversations between a reported violator and the reporter's account:

account: In a private Direct Message conversation, when a participant reports the other person, we will stop the violator from sending messages to the account who reported them. The conversation will also be removed from the reporter's inbox. However, if the reporter decides to continue to send Direct Messages to the violator, the conversation will resume.

Placing a Direct Message behind a notice: In a group Direct Message conversation, the violating Direct Message may be placed behind a notice to ensure no one else in the group can see it again.

Account-level enforcement

We take action at the account level if we determine that a person has violated the Twitter Rules in a particularly egregious way, or has repeatedly violated them even after receiving notifications from us.

Requiring media or profile edits: If an account's profile or media content is not compliant with our policies, we may make it temporarily unavailable and require that the violator edit the media or information in their profile to come into compliance. We also explain which policy their profile or media content has violated.

Placing an account in read-only mode: If it seems like an otherwise healthy account is in the middle of an abusive episode, we might temporarily make their account read-only, limiting their ability to Tweet, Retweet, or Like content until calmer heads prevail. The person can read their timelines and will only be able to send Direct Messages to their followers.

When an account is in read-only mode, others will still be able to see and engage with the account. The duration of this enforcement action can range from 12 hours to 7 days, depending on the nature of the violation.

Verifying account ownership: To ensure that violators do not abuse the anonymity we offer and harass others on the platform, we may require the account owner to verify ownership with a phone number or email address. This also helps us identify violators who are operating multiple accounts for abusive purposes and take action on such accounts. Note that when an account has been locked pending completion of a challenge (such as being required to provide a phone number), it is removed from follower counts, Retweets, and likes until it provides a phone number.

Permanent suspension: This is our most severe enforcement action. Permanently suspending an account will remove it from global view, and the violator will not be allowed to create new accounts. When we permanently suspend an account, we notify people that they have been suspended for abuse violations, and explain which policy or policies they have violated and which content was in violation.

Violators can appeal permanent suspensions if they believe we made an error. They can do this through the platform interface or by filing a report (<https://help.twitter.com/forms/general#abuse-suspended>). Upon appeal, if we find that a suspension is valid, we respond to the appeal with information on the policy that the account has violated.

Actions we may take against non-violating content

Placing a Tweet behind a notice: We may place some forms of [sensitive media] (<http://web.archive.org/en/rules-and-policies/media-policy.html>) like adult content or graphic violence behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. [This allows us to identify potentially sensitive content that some people may not wish to see.](#) Learn more about [how to control whether you see sensitive media] (<http://web.archive.org/en/safety-and-security/sensitive-media.html#settings>).

Withholding a Tweet based on age: We restrict views of specific forms of sensitive media such as adult content for viewers who are under 18 or viewers who do not include a birth date on their profile. [Adult content is any consensually produced and distributed media that is pornographic or intended to cause sexual arousal.](#) Age-restricted content will be placed behind interstitials, which you can learn about [here](https://help.twitter.com/en/rules-and-policies/notices-on-tweets) (<https://help.twitter.com/en/rules-and-policies/notices-on-tweets>).

Withholding a Tweet or account in a country: We may withhold access to certain content in a particular country if we receive a valid and properly scoped request from an authorized entity in that country. [We also clearly indicate within the product when content has been withheld.](#) Read more about [country withheld content] (<http://web.archive.org/en/rules-and-policies/tweet-withheld-by-country.html>).

In a private Direct Message conversation, when a participant reports the other account, we will stop the violator from sending messages to the account who reported them. The conversation will also be removed from the reporter's inbox. However, if the reporter decides to continue to send Direct Messages to the violator, the conversation will resume.

Placing a Direct Message behind a notice: In a group Direct Message conversation, the violating Direct Message may be placed behind a notice to ensure no one else in the group can see it again.

Account-level enforcement

We take action to suspend an account if we determine that a user has engaged in repeated violations of our policies and/or violated specific policies that cause significant risk to Twitter (i.e. posting illegal content, attempts to manipulate our platform or spam users, using our platform to incite violence, etc.) or pose a threat to our users (fraud, user privacy violations, violent threats, targeted harassment, etc.).

Placing an account in read-only mode: We may temporarily limit an account's ability to Tweet, Retweet, or Like. The user can read their timelines and will only be able to send Direct Messages to their followers.

Verifying account ownership: We may require an account owner to verify ownership with a phone number or email address. Note that when an account has been locked pending completion of a challenge (such as being required to provide a phone number), it is removed from follower counts, Retweets, and Likes.

Users can appeal account suspensions if they believe we made an error. They can do this through the platform interface or by [filing a report] (<http://web.archive.org/content/help-twitter/en/forms/account-access/appeals>).

Actions we may take against non-violating content

Placing a Tweet behind a notice: We may place some forms of [sensitive media] (<http://web.archive.org/en/rules-and-policies/media-policy.html>) like adult content or graphic violence behind an interstitial advising viewers to be aware that they will see sensitive media if they click through. [Learn more about \[how to control whether you see sensitive media\]](#) (<http://web.archive.org/en/safety-and-security/sensitive-media.html#settings>).

Withholding a Tweet based on age: We restrict views of specific forms of sensitive media such as adult content for viewers who are under 18 or viewers who do not include a birth date on their profile with [interstitials] (<http://web.archive.org/content/help-twitter/en/rules-and-policies/notices-on-tweets>).

Withholding a Tweet or account in a country: We may withhold access to certain content in a particular country if we receive a valid and properly scoped request from an authorized entity in that country. [Read more about \[country withheld content\]](#) (<http://web.archive.org/en/rules-and-policies/tweet-withheld-by-country.html>).

In a private Direct Message conversation, when a participant reports the other account, we will stop the violator from sending messages to the account who reported them. The conversation will also be removed from the reporter's inbox. However, if the reporter decides to continue to send Direct Messages to the violator, the conversation will resume.

Placing a Direct Message behind a notice: In a group Direct Message conversation, the violating Direct Message may be placed behind a notice to ensure no one else in the group can see it again.

Account-level enforcement

We take action to suspend an account if we determine that a user has engaged in repeated violations of our policies and/or violated specific policies that cause significant risk to ~~Twitter~~ (i.e. posting illegal content, attempts to manipulate our platform or spam users, using our platform to incite violence, etc.) or pose a threat to our users (fraud, user privacy violations, violent threats, targeted harassment, etc.).

Placing an account in read-only mode: We may temporarily limit an account's ability to ~~Tweet, Retweet,~~ or Like. The user can read their ~~timelines~~ and will only be able to send Direct Messages to their followers.

Verifying account ownership: We may require an account owner to verify ownership with a phone number or email address. Note that when an account has been locked pending completion of a challenge (such as being required to provide a phone number), it is removed from follower counts, ~~Retweets,~~ and Likes.

Users can appeal account suspensions if they believe we made an error. They can do this through the platform interface or by [filing a report] (<http://web.archive.org/content/help-twitter/en/forms/account-access/appeals>).

Actions we may take against non-violating content

Placing a ~~Tweet~~ behind a notice: We may place some forms of [sensitive ~~media~~] (<http://web.archive.org/en/rules-and-policies/media-policy.html>) like ~~adult content or graphic~~ violence behind an interstitial advising viewers to ~~be aware that they will~~ see sensitive media if they click through. [Learn more about \[how to control whether you see sensitive media\]](#) (<http://web.archive.org/en/safety-and-security/sensitive-media.html#settings>).

Withholding a ~~Tweet~~ based on age: We restrict views of specific forms of sensitive media such as adult content for viewers who are under 18 or viewers who do not include a birth date on their profile with [\[interstitials\]](#) (<http://web.archive.org/content/help-twitter/en/rules-and-policies/notices-on-tweets>).

Withholding a ~~Tweet~~ or account in a country: We may withhold access to certain content in a particular country if we receive a valid and properly scoped request from an authorized entity in that country. [Read more about \[country withheld content\]](#) (<http://web.archive.org/en/rules-and-policies/tweet-withheld-by-country.html>).

In a private Direct Message conversation, when a participant reports the other account, we will stop the violator from sending messages to the account who reported them. The conversation will also be removed from the reporter's inbox. However, if the reporter decides to continue to send Direct Messages to the violator, the conversation will resume.

Placing a Direct Message behind a notice: In a group Direct Message conversation, the violating Direct Message may be placed behind a notice to ensure no one else in the group can see it again.

Account-level enforcement

We take action to suspend an account if we determine that a user has engaged in repeated violations of our policies and/or violated specific policies that cause significant risk to X (i.e. posting illegal content, attempts to manipulate our platform or spam users, using our platform to incite violence, etc.) or pose a threat to our users (fraud, user privacy violations, violent threats, targeted harassment, etc.).

Placing an account in read-only mode: We may temporarily limit an account's ability to ~~post, Repost,~~ or Like. The user can read their ~~timelines~~ and will only be able to send Direct Messages to their followers.

Verifying account ownership: We may require an account owner to verify ownership with a phone number or email address. Note that when an account has been locked pending completion of a challenge (such as being required to provide a phone number), it is removed from follower counts, ~~Reposts,~~ and Likes.

Users can appeal account suspensions if they believe we made an error. They can do this through the platform interface or by [filing a report] (<http://en/forms/account-access/appeals.html>).

Actions we may take against non-violating content

Placing a ~~post~~ behind a notice: We may place some forms of [sensitive ~~media~~] (<http://en/rules-and-policies/media-policy.html>) like ~~adult content or graphic~~ violence behind an interstitial advising viewers to ~~be aware that they will~~ see sensitive media if they click through. [Learn more about \[how to control whether you see sensitive media\]](#) (<http://en/safety-and-security/sensitive-media.html>).

Withholding a ~~post~~ based on age: We restrict views of specific forms of sensitive media such as adult content for viewers who are under 18 or viewers who do not include a birth date on their profile with [\[interstitials\]](#) (<http://en/rules-and-policies/notices-on-x.html>).

Withholding a ~~post~~ or account in a country: We may withhold access to certain content in a particular country if we receive a valid and properly scoped request from an authorized entity in that country. [Read more about \[country withheld content\]](#) (<http://en/rules-and-policies/tweet-withheld-by-country.html>).