

Platform manipulation and spam policy

Overview

April 2022

You may not use **Twitter's** services in a manner intended to artificially amplify or suppress information or engage in behavior that manipulates or disrupts people's experience on **Twitter**.

We want **Twitter** to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using **Twitter** to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on **Twitter** to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting; **and**
- coordinated harmful activity that encourages or promotes behavior which violates [the \[Twitter Rules\]\(/en/rules-and-policies/twitter-rules.html\)](#).

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Multiple accounts and coordination

You can't mass-register **Twitter** accounts or use automation to create **Twitter** accounts.

Platform manipulation and spam policy

Overview

March 2023

You may not use **X 's** services in a manner intended to artificially **amplify** or suppress information or engage in behavior that manipulates or **disrupts** people's experience **or platform manipulation defenses on X**.

We want **X** to be a place where people can make human connections, find reliable information, and express themselves freely and safely. To make that possible, we do not allow spam or other types of platform manipulation. We define platform manipulation as using **X** to engage in bulk, aggressive, or deceptive activity that misleads others and/or disrupts their experience.

Platform manipulation can take many forms and our rules are intended to address a wide range of prohibited behavior, including:

- commercially-motivated spam, that typically aims to drive traffic or attention from a conversation on **X** to accounts, websites, products, services, or initiatives;
- inauthentic engagements, that attempt to make accounts or content appear more popular or active than they are;
- coordinated activity, that attempts to artificially influence conversations through the use of multiple accounts, fake accounts, automation and/or scripting;
- coordinated harmful activity that encourages or promotes behavior which violates [\[Our Rules\]\(http://web.archive.org/en/rules-and-policies/twitter-rules.html\)](#); and
- leveraging X's open source code to circumvent remediations or platform defenses.

What is in violation of this policy?

Under this policy we prohibit a range of behaviors in the following areas:

Multiple accounts and coordination

You can't mass-register **X** accounts or use automation to create **X** accounts.

X allows users to create and/or operate up to ten (10) accounts for different, non-duplicative purposes. Users who already own or operate more than ten accounts will not be impacted by these new limits unless a violation of the **X** Rules takes place. Some examples of what **X** considers to be non-duplicative uses include:

- Accounts tracking when objects in space travel over a specific location on Earth, e.g. New York, Hong Kong
- Accounts sharing news about different professional teams in the same sport or league, e.g. NFL, NBA, Premier League
- Accounts for personal projects, hobbies/interests, or for your business(es)
- Accounts for branded entities specific to unique locations or languages, e.g. brand accounts in English, Spanish and French

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate the **Twitter**

Rules. This includes:

- ~~operating multiple accounts with overlapping use cases, such as identical or similar personas or substantially similar content;~~
- ~~mutually interacting accounts—operating multiple accounts that interact with one another in order to inflate or manipulate the prominence of specific Tweets or accounts; and~~
- ~~coordination—creating multiple accounts to post duplicative content or create fake engagement, including:~~
 - ~~posting identical or substantially similar Tweets or hashtags from multiple accounts you operate;~~
 - ~~engaging (Retweets, Likes, mentions, Twitter Poll votes) repeatedly with the same Tweets or accounts from multiple accounts that you operate;~~
 - ~~coordinating with or compensating others to engage in artificial engagement or amplification, even if the people involved use only one account; and~~
 - ~~coordinating with others to engage in or promote violations of the Twitter Rules, including violations of our [abusive behavior policy]([en/rules-and-policies/abusive-behavior.html]).~~

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement.

This includes:

- selling/purchasing **Tweet** or account metric inflation – selling or purchasing followers or engagements (**Retweets**, Likes, mentions, **Twitter** Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to **Tweets**;
- reciprocal inflation – trading or coordinating to exchange follows or **Tweet** engagements (including but not limited to participation in “follow trains,” “decks,” and “**Retweet** for **Retweet**” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of **Twitter** accounts, usernames, or temporary access to **Twitter** accounts.

Misuse of **Twitter product features**

In addition, users are allowed to control multiple accounts on behalf of a third-party (e.g. “Social Media Managers”), provided that no violations of Our Rules (including this policy) occurs.

X may approve exceptions to multiple account limits at its discretion.

What is in violation of the Multiple Accounts policy?

You can't artificially amplify or disrupt conversations through the use of multiple accounts or by coordinating with others to violate Our Rules. This includes:

operating multiple accounts that interact with the same or substantially similar content or X account(s) in order to inflate or manipulate the prominence of content and/or accounts; for example:

- Do not create multiple accounts to “boost” trending topics or hashtags

- Do not create multiple accounts to engage with the same posts, accounts, or polls
- Do not create multiple accounts to misuse the mention/reply feature

- Do not create multiple accounts to boost or amplify one of your own accounts by misusing engagement features (likes, reposts)

operating multiple accounts that post substantially similar or identical content to one another; for example:

- Do not cross-post content across multiple accounts; content that is cross-posted but localized into another language is permitted

- Do not cross-post similar or duplicate content to the same trending topics or hashtags
- Employing workarounds to exceed the technical limits of account creation (e.g. [account phone number limits](<http://web.archive.org/en/managing-your-account/phone-number-faqs.html>))

Multiple accounts that exhibit a pattern or sole purpose of the preceding behaviors may be suspended by X. Users operating multiple accounts are still subject to all other X Rules, particularly our rules surrounding Platform Manipulation and Spam, and [Ban Evasion](<http://web.archive.org/en/rules-and-policies/ban-evasion.html>). Any other violations of Our Rules may also result in all of your accounts being suspended.

Engagement and metrics

You can't artificially inflate your own or others' followers or engagement.

This includes:

- selling/purchasing **Post** or account metric inflation – selling or purchasing followers or engagements (**Reposts**, Likes, mentions, **X** Poll votes);
- apps – using or promoting third-party services or apps that claim to add followers or add engagements to **Posts**;
- reciprocal inflation – trading or coordinating to exchange follows or **Post** engagements (including but not limited to participation in “follow trains,” “decks,” and “**Repost** for **Repost**” behavior); and
- account transfers or sales – selling, purchasing, trading, or offering the sale, purchase, or trade of **X** accounts, usernames, or temporary access to **X** accounts.

Misuse of X product features

You can't misuse **Twitter** product features to disrupt others' experience.

This includes:

Tweets and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical **Tweets**, or repeatedly sending identical Direct Messages;
- repeatedly posting **Tweets** or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your **Tweet/Direct** Message activity; and
- **Tweeting** an existing phrase or content in a duplicative manner, whether individually or in concert with other accounts. Learn more in our [copypasta and duplicate content **policy**](/en/rules-and-policies/copypasta-duplicate-content.html):

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with **Tweets** to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- **Tweeting** with excessive, unrelated hashtags in a single **Tweet** or across multiple **Tweets**:

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

You can't misuse **X** product features to disrupt others' experience.*• **This**

includes:

Posts and Direct Messages

- sending bulk, aggressive, high-volume unsolicited replies, mentions, or Direct Messages;
- posting and deleting the same content repeatedly;
- repeatedly posting identical or nearly identical **Posts**, or repeatedly sending identical Direct Messages;
- repeatedly posting **content** or sending Direct Messages consisting of links shared without commentary, so that this comprises the bulk of your **Post/Direct** Message activity; and
- **Posting** an existing phrase or content in a duplicative manner, whether individually or in concert with other accounts. Learn more in our [copypasta and duplicate content **policy**](http://web.archive.org/en/rules-and-policies/copypasta-duplicate-content.html).

Editing Posts

- editing **Posts** to artificially amplify content or deceive people. Examples include, but are not limited to:
- using a **Post's** existing engagement to amplify substantially different content (e.g. editing a **Post** from “What's better? Pancakes or waffles?” to “Thousands of people trust my service. Like my post and subscribe to my channel to get tips on investing”)
- editing links (**URLs**) so that the final destination page has significantly changed, either in content or location (e.g. domain, **URL path**)
- editing media (e.g. images, videos, audios, gifs, or **URLs** hosting relevant content) so the **Post** is meaningfully or entirely different or irrelevant to the previous versions.

Following

- “follow churn” – following and then unfollowing large numbers of accounts in an effort to inflate one's own follower count;
- indiscriminate following – following and/or unfollowing a large number of unrelated accounts in a short time period, particularly by automated means; and
- duplicating another account's followers, particularly using automation.

Engagement

- aggressively or automatically engaging with **Posts** to drive traffic or attention to accounts, websites, products, services, or initiatives.
- aggressively adding users to Lists or Moments.

Hashtags

- using a trending or popular hashtag with an intent to subvert or manipulate a conversation or to drive traffic or attention to accounts, websites, products, services, or initiatives; and
- **Posting** with excessive, unrelated hashtags in a single **Post** or across multiple **Posts**.

URLs

- publishing or linking to malicious content intended to damage or disrupt another person's browser (malware) or computer or to compromise a person's privacy (phishing); and
- posting misleading or deceptive links; e.g., affiliate links and clickjacking links.

What is not a violation of this policy?

The following are not in violation of this policy:

- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of ~~the Twitter~~ Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

~~Anyone can report accounts or Tweets via our dedicated reporting flow. These~~ reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Tweet*** from the icon.
2. Select **It 's suspicious or spam**.
3. Select the option that best tells us how the ~~Tweet~~ is suspicious or spreading spam.
4. Submit your report.

Desktop

You can report this content via ~~desktop~~ as follows:

1. Select **Report Tweet** from the icon.
2. Select **It 's suspicious or spam**.
3. Select the option that best tells us how the ~~Tweet~~ is suspicious or spreading spam.
4. Submit your report.

Report form

~~You can also report this content for review via our [spam reporting form](<https://help.twitter.com/forms/spam>) by selecting the I want to report spam on Twitter option.~~

Open Source Code

- exploiting X's open source code, including the logic behind visibility filtering or ML Model training, to circumvent policy enforcement for violations of [Our Rules] (<http://web.archive.org/content/help-twitter/en/rules-and-policies>).

What is not a violation of this policy?

The following are not in violation of this policy:

- posting links without commentary occasionally;
- coordinating with others to express ideas, viewpoints, support, or opposition towards a cause, provided such behavior does not result in violations of **Our** Rules; and
- operating multiple accounts with distinct identities, purposes, or use cases. These accounts may interact with one another, provided they don't violate other rules. Some examples include:
 - organizations with related but separate chapters or branches, such as a business with multiple locations;
 - operating a personal account in addition to pseudonymous accounts or accounts associated with your hobbies or initiatives; and
 - hobby/artistic bots.

Who can report violations of this policy?

Anyone can report accounts or Posts via our dedicated reporting flow. These reports are used in aggregate to help refine our enforcement systems and identify new and emerging trends and patterns of behavior.

How can I report violations of this policy?

In-app

You can report this content in-app as follows:

1. Select **Report Post*** from the icon.
2. Select **It 's suspicious or spam**.
3. Select the option that best tells us how the **Post** is suspicious or spreading spam.
4. Submit your report.

Web

You can report this content via **web** as follows:

1. Select **Report Post** from the icon.
2. Select **It 's suspicious or spam**.
3. Select the option that best tells us how the **Post** is suspicious or spreading spam.
4. Submit your report.

What happens if you violate this policy?

What happens if you violate this policy?

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [\[unsafe links\]\(/en/safety-and-security/phishing-spam-and-malware-links.html\)](#), including how to appeal if we've falsely [identified your URL as unsafe](#).

Tweet deletion and temporary account locks

• If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one of more Tweets to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.

• In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that [Twitter](#) is able to reliably attribute to entities known to violate [the \[Twitter Rules\]\(/en/rules-and-policies/twitter-rules.html\)](#).

The consequences for violating this policy depend on the severity of the violation as well as any previous history of violations. Our action is also informed by the type of spammy activity that we have identified. The actions we take may include the following:

Anti-spam challenges

When we detect suspicious levels of activity, accounts may be locked and prompted to provide additional information (e.g., a phone number) or to solve a reCAPTCHA.

Denylisting URLs

We denylist or provide warnings about URLs we believe to be unsafe. Read more about [\[unsafe links\]\(http://web.archive.org/en/safety-and-security/phishing-spam-and-malware-links.html\)](#), including how to appeal if we've falsely [identified your URL as unsafe](#).

Limiting the visibility of Posts

Posts which we believe to be in violation of these policies may not appear in

certain parts of the X product, and/or may not be recommended or amplified by X. [\[Learn more\]\(http://web.archive.org/en/rules-and-policies/twitter-reach-limited.html\)](#) about other instances when a Post's visibility may be limited.

Post deletion and temporary account locks

• If the platform manipulation or spam offense is an isolated incident or first offense, we may take a number of actions ranging from requiring deletion of one of more Posts to temporarily locking account(s). Any subsequent platform manipulation offenses will result in permanent suspension.

• In the case of a violation centering around the use of multiple accounts, you may be asked to choose one account to keep. The remaining accounts will be permanently suspended.

Temporary loss of access to X features or products

• We may temporarily limit or restrict access to X features or products, including (but not limited to) Posts, Edit Post, Direct Messages, Spaces, Communities, or Live.

Permanent suspension

For severe violations, accounts will be permanently suspended at first detection. Examples of severe violations include:

- operating accounts where the majority of behavior is in violation of the policies described above;
- using any of the tactics described on this page to undermine the integrity of elections;
- buying/selling accounts;
- creating accounts to replace or mimic a suspended account; and
- operating accounts that [X](#) is able to reliably attribute to entities known to violate [\[Our Rules\]\(http://web.archive.org/en/rules-and-policies/twitter-rules.html\)](#).

If you believe that your account was locked or suspended in error, you can [submit an appeal](<https://help.twitter.com/forms/general?subtopic=suspended>).

Additional resources

Learn more about our [automation rules](<https://about.twitter.com/en/rules-and-policies/twitter-automation.html>) for developers, our [election integrity](<https://about.twitter.com/en/values/elections-integrity.html>) efforts, our [financial scam policy](<https://about.twitter.com/en/rules-and-policies/financial-scam.html>), our [hacked materials policy](<https://about.twitter.com/en/rules-and-policies/hacked-materials.html>), our approach to [coordinated harmful activity](<https://about.twitter.com/en/rules-and-policies/coordinated-harmful-activity.html>), and our [guidelines for promotions and contests](<https://about.twitter.com/en/rules-and-policies/twitter-contest-rules.html>).

Learn more about our [range of enforcement options](<https://about.twitter.com/en/rules-and-policies/enforcement-options.html>) and our approach to [policy development and enforcement](<https://about.twitter.com/en/rules-and-policies/enforcement-philosophy.html>).

If you believe that your account was locked or suspended in error, you can [submit an appeal](<https://help.twitter.com/forms/general?subtopic=suspended>).

Additional resources

Learn more about our [automation rules](<http://web.archive.org/en/rules-and-policies/twitter-automation.html>) for developers, our [election integrity](<https://about.twitter.com/en/values/elections-integrity.html>) efforts, our [financial scam policy](<http://web.archive.org/en/rules-and-policies/financial-scam.html>), our [hacked materials policy](<http://web.archive.org/en/rules-and-policies/hacked-materials.html>), and our [guidelines for promotions and contests](<http://web.archive.org/en/rules-and-policies/twitter-contest-rules.html>).

Learn more about our [range of enforcement options](<http://web.archive.org/en/rules-and-policies/enforcement-options.html>) and our approach to [policy development and enforcement](<http://web.archive.org/en/rules-and-policies/enforcement-philosophy.html>).